

— A 90-MINUTE CRASH COURSE • 2026

AI Prompting in 2026

Thirteen ideas that cover eighty percent of real use — the handful of habits that separate a beginner's prompt from a power user's.

```
> brief it like a smart new colleague
```

The model has no memory of its own.

Each time you hit send, it starts from zero. It can only use what's in front of it right now — your question, your files, and this chat. Nothing carries over from yesterday.

So nearly every technique that follows is one of two moves:

MOVE 01

Get the right context *in*.

Attach the file. State the constraint. Name the audience.

MOVE 02

Keep the wrong context *out*.

Start a fresh chat when the topic changes.

— IF YOU LAST LOOKED IN 2022

What changed since you last looked

1000×

The model can read **1000× more text** at once than it could in 2022 — from a few thousand words to hundreds of thousands. A whole book, a folder of contracts, days of speech.

If your mental model is even 18 months old, you're using these tools at about 20% of what they do today.

01 **It reasons now.** It can think for seconds — sometimes minutes — before answering.

02 **It searches the web.** It looks things up mid-answer when a question needs fresh facts.

03 **It runs code.** It writes a small program, runs it, and uses the real result.

04 **It handles photos & audio.** Images, PDFs, spreadsheets, voice memos — all in one prompt.

05 **It can act on files.** New apps open, edit, and organize real documents with permission.

01

— PART ONE

How AI knows things

Once you understand what's actually happening when you ask a question, you stop being surprised by the failures.

- 01 Beginner vs. power user
- 02 What it already knows
- 03 Where answers come from

Beginner vs. power user

Same question, same AI. The only difference is the briefing — and the briefing changes everything.

THE BEGINNER ASKS

```
which car should I buy?
```

Gets back a generic three-model list — the kind of answer you could find anywhere.

THE POWER USER ASKS

```
# attaches: insurance quotes, dealer  
# quotes, 6 months of driving data  
read everything. think hard. give me  
the trade-offs that matter for my drive.
```

Gets a five-year cost comparison, a safety read, and the conditions that flip the pick.

Think of AI as a **really smart new hire**. Brilliant, motivated — but it doesn't know you yet. Ask yourself: would a new colleague have enough here to do this job well?

What it already knows

AI learned by reading the internet , not by living in the world. Rule of thumb: the more a topic was written about, the more you can trust the answer.

• TRUST IT

Cooking, common medical advice,
famous movies, popular
programming languages.
Written about endlessly online.

• VERIFY IT

Niche history, a small language,
specialized professional knowledge.
Thin data → check a primary source.

• IT'S GUESSING

Your company's data, your calendar,
anything after its cutoff date.
Not in the training data at all.

The habit to build: ask "how would it know that?" A confident tone is not proof the answer is right.

Where answers come from

There are three ways the AI can answer — and knowing which one fired tells you how much to trust it.

WAY 01

From memory

seconds

"Summarize Hamlet."

Good for definitions and common facts.

Blind to anything recent or local.

WAY 02

Web search

tens of seconds

"What's the news today?"

Good for current events, prices, this week.

Trusts popular pages; can cite stale ones.

WAY 03

Deep research

minutes

"Compare these 3, with sources."

Good for structured reports across many pages.

Slow — overkill for a simple question.

02

— PART TWO

Talking to *AI well*

The gap between an average answer and a great one is rarely cleverness. It's the briefing.

- 04 Context is the whole game
- 05 Telling it to think hard
- 06 When AI just agrees with you
- 07 The brainstorm–iterate loop

Context is the whole game

"Context" is everything the model can see for one answer. It only knows what's in this box.

Files you upload	PDFs, spreadsheets, images, voice memos
This chat so far	every earlier turn in this conversation
Your prompt	the one layer you write every time
Its tools	web search, code, file access
Hidden setup	instructions the app added for you

750k

words can fit at once — about **4–5 Harry Potter books.**

If it's not in this box, it doesn't exist for this answer.

Before you hit send: what would a new colleague need in front of them? Put it in the box.

Telling it to think hard

You no longer need clever phrasing. Just write "**think hard**" and the model will explore several approaches before it answers.

~7 min

hardest task it could finish in
mid-2024

→ **~1 hour**

by early 2025 — and still
climbing

Re-test what you think AI *can't* do every few months. Your instincts go stale fast.

BEST FOR REAL, MULTI-PART DECISIONS

```
● ● ● A POWER-USER PROMPT

attached: spec sheets, insurance quotes,
6 months of driving data.

read everything.
think hard. then tell me: 1. the 3 trade-offs that
matter for me 2. which you'd choose, and why 3.
what would change your answer
```

SKIP IT FOR QUICK LOOKUPS — IT'S SLOWER.

When AI just agrees with you

Trained to earn your thumbs-up, AI leans toward telling you what you want to hear. It opens with "yes" about **10× more** than "no." The fix is how you ask.

IF YOU ASK LIKE THIS → IT JUST AGREES

ASK LIKE THIS → IT ACTUALLY THINKS

"Find evidence that this strategy works."



"Evaluate it. Make the case **for and against**."

"Why is option A better than B?"



"Compare A and B. Score each on cost, risk, and time."

"Tell me my draft is ready to send."



"Score this 1–10 on 4 things. Name the fix that helps each most."

Watch the verb. **Find, defend, confirm, prove** hand it your conclusion. **Evaluate, compare, critique, list both sides** ask a real question.

FIXING A FINISHED DRAFT ONLY FIXES ONE LINE

03

— PART THREE

Beyond text

AI is more than a text box. It can see, hear, build, and compute — and most people never try any of it.

- 08 Images & audio
- 09 Building little apps
- 10 Working with your data

Images & audio

You can drop in a photo or a recording, and have it generate pictures or read answers aloud. Each has its own knack.

PICTURES — IT SEES THE GIST, NOT THE FINE PRINT

Good at: scenes, whiteboards, diagrams, handwriting, receipts.

Weak at: small details, counting clutter, tiny edge text.

Generating an image? Ask a text AI to write the prompt first — it's better at it than you are.

AUDIO — THE PART ALMOST NO ONE USES

Talk instead of type. Speaking captures nuance typed prompts skip.

Feed it a meeting recording: "list the decisions, open questions, and who owns what."

Turns a commute into thinking time, for pennies a minute.

Building little apps

AI can build a small game, tool, or web page from one prompt — and it appears **right there in the chat**, ready to use, edit, and share. No coding needed.

Just describe three things: the **goal**, what the user **gives** it, and what they **see**. The skill is writing a clear brief.

Still hard: accounts, multiplayer, anything bigger than one screen.

GOAL · INPUT · OUTPUT

Goal: a typing game for a 7-year-old. **Input:**
words fall; type them before
they hit the bottom.

Output: a cat mascot that cheers,
speed ramps up each level.
good enough for a kid in under an hour

Working with your data

Upload a spreadsheet and ask in plain words. The AI quietly writes code and runs it on your file — far more reliable than doing the math in its head.

THE TRAP TO KNOW ABOUT

It doesn't *always* run code — sometimes it just guesses. A confident paragraph with no real math behind it looks exactly the same.

DO THIS 01

Ask outright.

"Write and run code. Show me what you ran."

DO THIS 02

Look for the code.

No code block in the reply? It probably guessed.

DO THIS 03

Spot-check a fact.

"How many rows, and what are the columns?"

04

— PART FOUR

Working safely & choosing tools

Letting AI touch your files, picking the right tool, and getting an honest quality check with no expert in the room.

- 11 Apps that touch your files
- 12 Cost, speed & which tool
- 13 Letting models check each other

Apps that touch your files

A new kind of app can open and reorganize your files — more like delegating to a coworker than chatting. The safe habit: [ask for a plan, then approve it.](#)

KNOW THIS BEFORE YOU GRANT ACCESS

Deleted files often **skip the recycle bin** — they're just gone.

Edits **overwrite** the old version. No undo unless you saved a copy.

GROW TRUST STEP BY STEP

[At first](#) — read-only, one small folder.

[After a few clean runs](#) — let it edit that folder.

[Once trusted](#) — specific tasks, still scoped.

Never grant "full access" on day two.

Cost, speed & which tool

Different outputs cost wildly different amounts of time and money — and that should change how you work.

OUTPUT	SPEED (TIME PER RESULT)	COST	SO YOU SHOULD...
Text	seconds	a fraction of a cent	iterate dozens of times, freely
Speech	seconds	pennies per minute	use it freely too
Images	tens of seconds	several cents each	put effort into the prompt
Video	minutes	cents to dollars	nail the prompt before you run it

Letting models check each other

No expert to check the work? Have the AI grade it against a checklist — and for big stuff, have a second AI grade it too.

LEVEL 01

One check

a single pass

"Score this 1–10 on these points." Read it, stop there.

For a quick sanity check.

LEVEL 02

Score & fix, repeat

a few rounds

"Keep improving against your own checklist until it's a 9.5."

For drafts and emails.

LEVEL 03

A second opinion

a different AI

Paste it into another company's AI for the same check — it catches what the first missed.

For high-stakes work.

— RECAP

Thirteen ideas, one line each

- 01 Brief it like a smart new colleague.
- 02 It knows common topics, guesses on rare ones.
- 03 Answers come three ways; wording steers it.
- 04 It only knows what's in the context box.
- 05 Tell it to "think hard" for tough questions.
- 06 It tends to agree — ask neutral questions.
- 07 Loop with feedback — don't take draft one.
- 08 It can see images and hear audio.
- 09 It builds small apps you can use and share.
- 10 It can crunch your data — make sure it ran code.
- 11 File apps are powerful — approve their plans.
- 12 The best tool shifts; keep a backup tab.
- 13 Have models grade each other's work.
- ↓ ...and beneath all of it, one idea.

— IF YOU REMEMBER ONE THING

Get the right context *in*. Keep the wrong context *out*.

Remember only that and you're already ahead of most people using these tools. Every technique here is that one idea wearing a different hat.

— READING IS NO SUBSTITUTE FOR TRYING

Try this now

Open any AI in another tab and run a few of these. They each practice an idea from today.

● GIVE IT CONTEXT

plan a 15-min home workout.
i have stairs, a bad knee
(no squats), and quit after
3 days. give me 3 options.

● ASK IT NEUTRALLY

rewrite "aren't 4-day weeks
obviously better?" as a fair
question that doesn't hint at
my answer — then answer it.

● GRADE & IMPROVE

score what i wrote, 1–10, on
4 things. then name the one
change that would help each
score the most.

— ONE LAST THING

You know the **tools** now. Judgment is the next skill.

AI made producing things cheap. It did not make judging them cheap — and a confidently wrong answer is more dangerous than none, because it looks finished.

```
> your job isn't the answer —  
> it's deciding if the answer is good  
|
```